

# Lecture 17: Multitasking

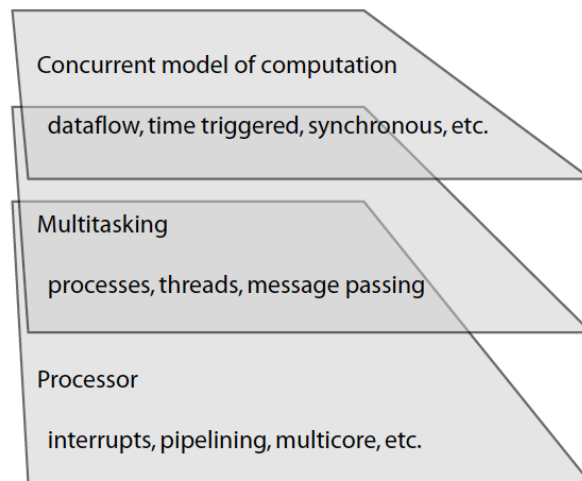
Seyed-Hosein Attarzadeh-Niaki

Based on the slides by Edward Lee

Real-Time Embedded Systems

1

## Layers of Abstraction for Concurrency in Programs



Real-Time Embedded Systems

2

## Definition and Uses

**Threads** are sequential procedures that share memory.

Uses of **concurrency**?

- Reacting to external events (interrupts)
- Exception handling (software interrupts)
- Creating the illusion of simultaneously running different programs (multitasking)
- Exploiting parallelism in the hardware (e.g. multicore machines).
- Dealing with real-time constraints.

## Thread Scheduling

Predicting the thread schedule is an iffy proposition.

- Without an OS, multithreading is achieved with **interrupts**. Timing is determined by external events.
- Generic OSs (Linux, Windows, OSX, ...) provide **thread libraries** (like “pthreads”) and provide no fixed guarantees about when threads will execute.
- **Real-time operating systems (RTOSs)**, like FreeRTOS, QNX, VxWorks, RTLinux, support a variety of ways of controlling when threads execute (priorities, preemption policies, deadlines, ...).
- Processes are collections of threads with their own memory, not visible to other processes. *Segmentation faults* are attempts to access memory not allocated to the process. Communication between processes must occur via OS facilities (like pipes or files).

## Posix Threads (PThreads)

- PThreads is an API (Application Program Interface) implemented by many operating systems, both real-time and not.
  - It is *a library of C procedures*.
- Standardized by the IEEE in 1988 to unify variants of Unix.
  - Subsequently implemented in most other operating systems.
- An alternative is Java, which may use PThreads under the hood, but provides thread constructs as part of the programming language.

Real-Time Embedded Systems

5

## Creating and Destroying Threads

```
#include <pthread.h>

void* threadFunction(void* arg) {
    ...
    return pointerToSomething or NULL;
}

int main(void) {
    pthread_t threadID;
    void* exitStatus;
    int value = something;
    pthread_create(&threadID, NULL, threadFunction, &value);
    ...
    pthread_join(threadID, &exitStatus);
    return 0;
}
```

Can pass in pointers to shared variables.

Can return pointer to something.  
Do not return a pointer to a local variable!

Create a thread (may or may not start running!)

Becomes arg parameter to threadFunction.  
Why is it OK that this is a local variable?

Return only after all threads have terminated.

Real-Time Embedded Systems

6

## What's Wrong with This?

```
#include <pthread.h>
#include <stdio.h>
void *myThread() {
    int ret = 42;
    return &ret;
}
```

Don't return a pointer to a local variable, which is on the stack.

```
int main() {
    pthread_t tid;
    void *status;
    pthread_create(&tid, NULL, myThread, NULL);
    pthread_join(tid, &status);
    printf("%d\n", *(int*)status); return 0;
}
```

Real-Time Embedded Systems

7

## Notes

- Threads can (and often do) share variables
- Threads may or may not begin running immediately after being created.
- A thread may be suspended between any two atomic instructions (typically, assembly instructions, not C statements!) to execute another thread and/or interrupt service routine.
- Threads can often be given priorities, and these may or may not be respected by the thread scheduler.
- Threads may block on semaphores and mutexes (we will do this later in this lecture).

Real-Time Embedded Systems

8

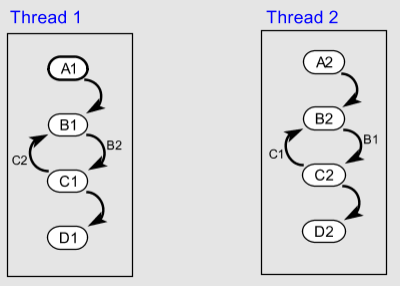
# Modeling Threads via Asynchronous Composition of Extended State Machines

States or transitions represent atomic instructions

## Interleaving semantics

- Choose one machine, arbitrarily.
- Advance to a next state if guards are satisfied.
- Repeat.

Need to compute reachable states to reason about correctness of the composed system



Can Thread 1 be in C1 at the same time Thread 2 is in C2?

Real-Time Embedded Systems

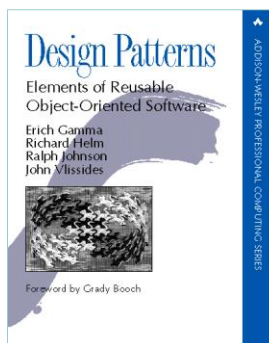
9

## A Scenario

Under Integrated Modular Avionics, software in the aircraft engine continually runs diagnostics and publishes diagnostic data on the local network.



Proper software engineering practice suggests using the observer pattern.



An observer process updates the cockpit display based on notifications from the engine diagnostics.

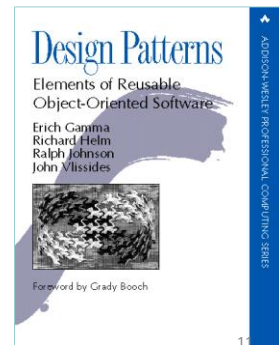
Real-Time Embedded Systems

10

## Typical thread programming problem

“The *Observer pattern* defines a one-to-many dependency between a subject object and any number of observer objects so that when the subject object changes state, all its observer objects are notified and updated automatically.”

*Design Patterns*, Eric Gamma, Richard Helm, Ralph Johnson, John Vlissides (Addison-Wesley, 1995)



Real-Time Embedded Systems

1

## Observer Pattern in C

```
// Value that when updated triggers notification of
// registered listeners.
int value;

// List of listeners. A linked list containing
// pointers to notify procedures.
typedef void* notifyProcedure(int);
struct element {...}
typedef struct element elementType;
elementType* head = 0;
elementType* tail = 0;

// Procedure to add a listener to the list.
void addListener(notifyProcedure listener) {...}

// Procedure to update the value
void update(int newValue) {...}

// Procedure to call when notifying
void print(int newValue) {...}
```

Real-Time Embedded Systems

12

# Observer Pattern in C

```
// Value that when updated triggers notification of
// registered listeners.
int value;

// List of listeners. A list of pointers to notify procedure.
typedef void* notifyProcedure;
struct element {
    notifyProcedure* listener;
    struct element* next;
};
typedef struct element elementType;
elementType* head = 0;
elementType* tail = 0;

// Procedure to add a listener to the list.
void addListener(notifyProcedure listener) {...}

// Procedure to update the value
void update(int newValue) {...}

// Procedure to call when notifying
void print(int newValue) {...}
```

Real-Time Embedded Systems

13

# Observer Pattern in C

```
// Value that when updated triggers notification of
// registered listeners.
int value;

// List of listeners. A list of pointers to notify procedure.
typedef void* notifyProcedure;
struct element {
    notifyProcedure* listener;
    struct element* next;
};
typedef struct element elementType;
elementType* head = 0;
elementType* tail = 0;

// Procedure to add a listener to the list.
void addListener(notifyProcedure listener) {
    if (head == 0) {
        head = malloc(sizeof(elementType));
        head->listener = listener;
        head->next = 0;
        tail = head;
    } else {
        tail->next = malloc(sizeof(elementType));
        tail = tail->next;
        tail->listener = listener;
        tail->next = 0;
    }
}

// Procedure to update the value
void update(int newValue) {...}

// Procedure to call when notifying
void print(int newValue) {...}
```

Real-Time Embedded Systems

14

# Observer Pattern in C

```
// Value that when updated triggers notification of
// registered listeners.
int value;

// List of listeners. A linked list containing
// pointers to notify procedures.
typedef void* notifyProcedure(int);
struct element {...}
typedef struct element* element_type;
element_type head;
element_type element;

// Procedure to update the value
void update(int newValue) {
    value = newValue;
    // Notify listeners.
    element_type element = head;
    while (element != 0) {
        (*(element->listener))(newValue);
        element = element->next;
    }
}

// Procedure to add a listener
void addListener(notifyProcedure listener) {
    // ...
}

// Procedure to print the value
void print(int value) {
    // ...
}
```

Real-Time Embedded Systems

15

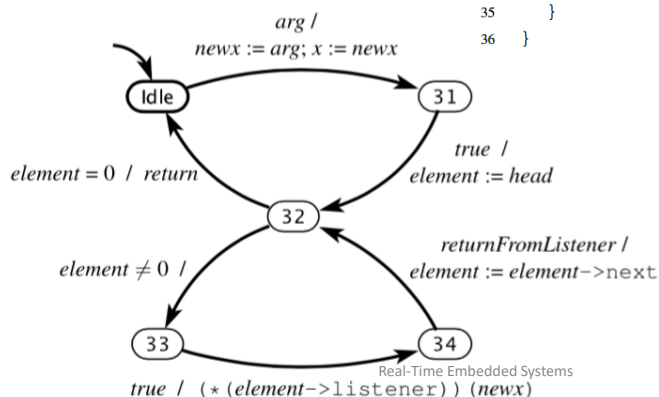
## Model of the Update Procedure

**inputs:** *arg*: int, *returnFromListener*: pure

**outputs:** *return*: pure

**local variables:** *newx*: int, *element*: element\_t\*

**global variables:** *x*: int, *head*: element\_t\*



Real-Time Embedded Systems

16



# Observer Pattern in C

```
// Value that when updated triggers notification of registered listeners.
int value;

// List of listeners. A linked list containing
// pointers to notify procedures.
typedef void* notifyProcedure(int);
struct element {...}
typedef struct element elementType;
elementType* head = 0;
elementType* tail = 0;

// Procedure to add a listener to the list.
void addListener(notifyProcedure listener) {...}

// Procedure to update the value
void update(int newValue) {...}

// Procedure to call when notifying
void print(int newValue) {...}
```

Will this work in a  
multithreaded context?

Will there be  
unexpected/undesirable  
behaviors?

What if addListener is  
called from two threads at  
the same time?

Real-Time Embedded Systems

17

```
#include <pthread.h>
...
pthread_mutex_t lock;

void addListener(notify listener) {
    pthread_mutex_lock(&lock);
    ...
    pthread_mutex_unlock(&lock);
}

void update(int newValue) {
    pthread_mutex_lock(&lock);
    value = newValue;
    elementType* element = head;
    while (element != 0) {
        (*(element->listener))(newValue);
        element = element->next;
    }
    pthread_mutex_unlock(&lock);
}

int main(void) {
    pthread_mutex_init(&lock, NULL);
    ...
}
```

## Using Posix mutexes on the observer pattern in C

However, this carries a  
significant **deadlock risk**.  
The update procedure  
holds the lock while it  
calls the notify  
procedures. If any of  
those stalls trying to  
acquire another lock, and  
the thread holding that  
lock tries to acquire this  
lock, deadlock results.

Real-Time Embedded Systems

18

## One possible “fix”

```
#include <pthread.h>
...
pthread_mutex_t lock;

void addListener(notify listener) {
    pthread_mutex_lock(&lock);
    ...
    pthread_mutex_unlock(&lock);
}

void update(int newValue) {
    pthread_mutex_lock(&lock);
    value = newValue;
    ... copy the list of listeners ...
    pthread_mutex_unlock(&lock);
    elementType* element = headCopy;
    while (element != 0) {
        (*(element->listener))(newValue);
        element = element->next;
    }
}

int main(void) {
    pthread_mutex_init(&lock, NULL);
    ...
}
```

What is wrong with this?

Notice that if multiple threads call update(), the updates will occur in some order. But there is no assurance that the listeners will be notified in the same order. Listeners may be misled about the “final” value.

Real-Time Embedded Systems

19

This is a very simple, commonly used design pattern. Perhaps Concurrency is Just Hard...

Sutter and Larus observe:

*“Humans are quickly overwhelmed by concurrency and find it much more difficult to reason about concurrent than sequential code. Even careful people miss possible interleavings among even simple collections of partially ordered operations.”*

*H. Sutter and J. Larus. Software and the concurrency revolution. ACM Queue, 3(7), 2005.*

Real-Time Embedded Systems

20

If concurrency were intrinsically hard, we would not function well in the physical world



*It is not  
concurrency that  
is hard...*

Real-Time Embedded Systems

21

...It is Threads that are Hard!

Threads are sequential processes that share memory. From the perspective of any thread, *the entire state of the universe can change between any two atomic actions* (itself an ill-defined concept).

*Imagine if the physical world did that...*

Real-Time Embedded Systems

22

# What it Feels Like to Use Mutexes



Image "borrowed" from an Omega advertisement for Y2K software and disk drives, Scientific American, September 1999.

Real-Time Embedded Systems

23

## Message-passing programs *may* be better

```

1 void* producer(void* arg) {
2     int i;
3     for (i = 0; i < 10; i++) {
4         send(i);
5     }
6     return NULL;
7 }
8 void* consumer(void* arg) {
9     while(1) {
10         printf("received %d\n", get());
11     }
12     return NULL;
13 }
14 int main(void) {
15     pthread_t threadID1, threadID2;
16     void* exitStatus;
17     pthread_create(&threadID1, NULL, producer, NULL);
18     pthread_create(&threadID2, NULL, consumer, NULL);
19     pthread_join(threadID1, &exitStatus);
20     pthread_join(threadID2, &exitStatus);
21     return 0;
22 }
```

But there is still risk of  
deadlock and  
unexpected  
nondeterminism!

Real-Time Embedded Systems

24

## Claim

*Nontrivial software written with threads, semaphores, and mutexes is incomprehensible to humans.*

- *Need better ways to program concurrent systems (we will see some later in the course)*
- *Better tools to analyze and reason about concurrency (e.g. model checking)*

## Do Threads Have a Sound Foundation?

If the foundation is bad, then we either tolerate brittle designs that are difficult to make work, or we have to rebuild from the foundations.

Note that this whole enterprise is held up by threads



# Problems with the Foundations

A model of computation:

Bits:  $B = \{0, 1\}$

Set of finite sequences of bits:  $B^*$

Computation:  $f: B^* \rightarrow B^*$

Composition of computations:  $f \bullet f'$

Programs specify compositions of computations

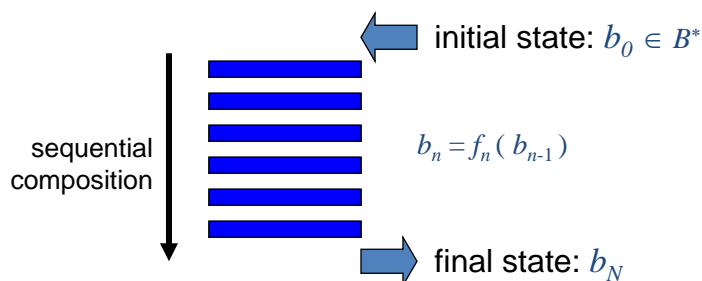
*Threads augment this model to admit concurrency.*

*But this model does not admit concurrency gracefully.*

Real-Time Embedded Systems

27

# Basic Sequential Computation

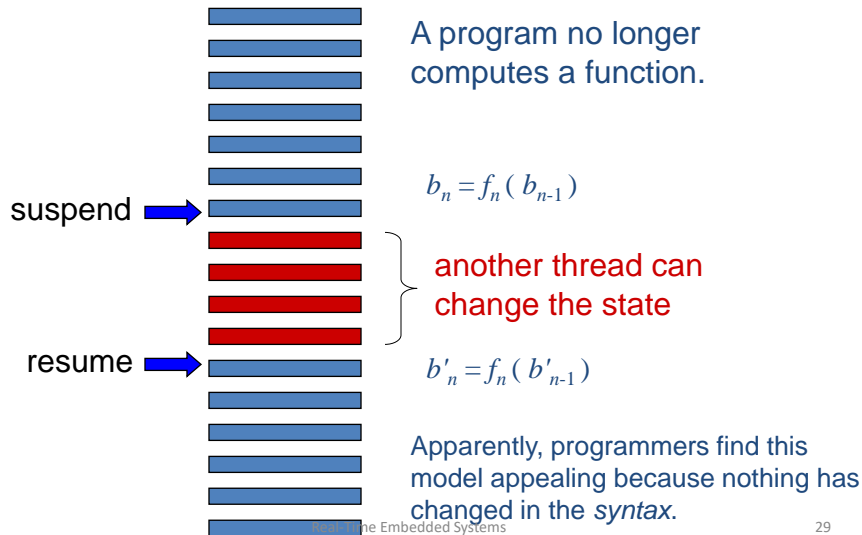


*Formally, composition of computations is function composition.*

Real-Time Embedded Systems

28

## When There are Threads, Everything Changes



## Succinct Problem Statement

*Threads are wildly **nondeterministic**.*

The programmer's job is to prune away the nondeterminism by imposing constraints on execution order (e.g., mutexes) and limiting shared data accesses (e.g., OO design).

## Incremental Improvements to Threads

- Object Oriented programming
- Coding rules (Acquire locks in the same order...)
- Libraries (Stapl, Java  $\geq 5.0$ , ...)
- Transactions (Databases, ...)
- Patterns (MapReduce, ...)
- Formal verification (Model checking, ...)
- Enhanced languages (Split-C, Cilk, Guava, ...)
- Enhanced mechanisms (Promises, futures, asynchronous atomic callbacks ...)

## *IEEE Computer*, May, 2006

