

Lecture 21: Real-Time Operating Systems

Seyed-Hosein Attarzadeh-Niaki

Some slides due to Ingo Sander and Edward Lee

Embedded Real-Time Systems

1

Review

- Mutual exclusion
 - Priority inversion
 - Priority inheritance protocol
 - Deadlock
 - Priority ceiling protocol
- Aperiodic scheduling
 - Polling server
 - Sporadic server
- Multiprocessor scheduling
 - Brittleness
 - Richard's anomalies


Embedded Real-Time Systems

2


Disk and Network Handled by Tasks

- Effectively no device needs to be supported by all variants of the OS, except maybe the system timer.
- Disk & network handled by tasks instead of integrated drivers

Embedded OS



Standard OS



Embedded Real-Time Systems

5

Protection is Optional

- **Protection** mechanisms **not always necessary**
 - ES typically designed for a single purpose
 - Untested programs rarely loaded
 - SW considered reliable.
- **Privileged I/O** instructions **not necessary** and tasks can do their own I/O.

Embedded Real-Time Systems

6

Interrupts not restricted to OS

- **Interrupts** can be employed by **any process**
 - For standard OS: serious source of unreliability.
 - For embedded systems: efficient control over a variety of devices is required
- Embedded programs can be considered to be **tested**
- **More efficient** than going through OS services

Real-time capability


- The timing behavior of the OS must be **predictable**
 - \forall services of the OS: **Upper bound on the execution time!**
- OS should manage the **timing** and **scheduling**
 - *Internal synchronization*: with one master clock or distributed
 - *External synchronization*: guarantees consistency with actual physical time (GPS, etc.)
- The OS must be *fast*
 - Practically important

Functionality of RTOS-Kernels

- processor management
- memory management,
- and timer management; } resource management
- task management (resume, wait etc),
- inter-task communication and synchronization.

Classes of RTOSs

- Fast proprietary kernels
 - QNX, PDOS, VxWorks, VxWORKS
- RT extensions to standard Oss



- Research RTOSs trying to avoid limitations
 - MARS, Spring, MARUTI, Arts, Hartos, DARK, and Melody

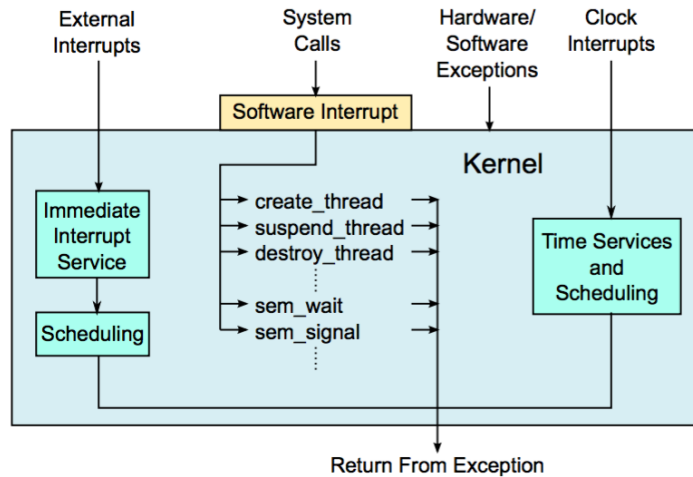
Outline of a Microkernel Scheduler

- Scheduler may be part of a compiler or code generator, part of an OS or microkernel, or both.
- Main Scheduler Thread (Task)
 - set up periodic timer interrupts;
 - create default thread data structures;
 - dispatch a thread (procedure call);
 - execute main thread (idle or power save, for example).
- Thread data structure
 - copy of all state (machine registers)
 - address at which to resume executing the thread
 - status of the thread (e.g. blocked on mutex)
 - priority, WCET (worst case execution time), and other info to assist the scheduler

Outline of a Microkernel Scheduler

- Timer interrupt service routine:
 - dispatch a thread.
- Upon dispatching a thread
 - *disable interrupts;*
 - determine which thread should execute (scheduling);
 - if the same one, enable interrupts and return;
 - save state (registers) into current thread data structure;
 - save return address from the stack for current thread;
 - copy new thread state into machine registers;
 - replace program counter on the stack for the new thread;
 - *enable interrupts;*
 - return.

Structure of A Microkernel

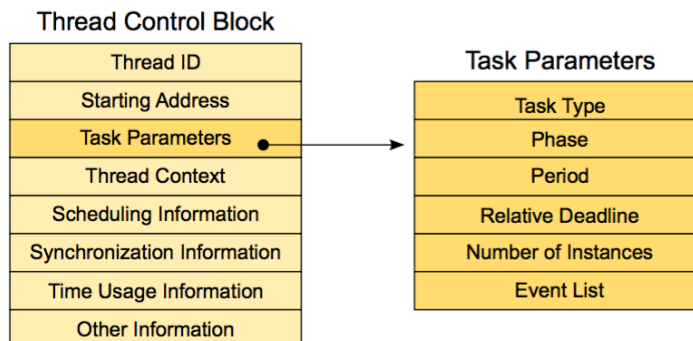


Embedded Real-Time Systems

13

Thread Control Block (TCB)

The Thread Control Block keeps all information needed to manage and schedule the thread.



Embedded Real-Time Systems

14

Thread Control Block (TCB)

When a thread is

- *executing*, its context is changed continuously.
- *stops execution*, its context is stored in the thread's TCB.
- *'put in a queue'* (ready queue, pending queue for semaphore, . . .), a pointer to the threads TCB is put in the linked list representing the queue.
- *destroyed*, the kernel deletes the TCB and deallocates its memory space.

Task (Implementing) Function

```
void ATaskFunction( void *pvParameters )
{
    /* Variables can be declared just as per a normal function. Each instance
    of a task created using this function will have its own copy of the
    iVariableExample variable. This would not be true if the variable was
    declared static - in which case only one copy of the variable would exist
    and this copy would be shared by each created instance of the task. */
    int iVariableExample = 0;

    /* A task will normally be implemented as an infinite loop. */
    for( ;; )
    {
        /* The code to implement the task functionality will go here. */
    }

    /* Should the task implementation ever break out of the above loop
    then the task must be deleted before reaching the end of this function.
    The NULL parameter passed to the vTaskDelete() function indicates that
    the task to be deleted is the calling (this) task. */
    vTaskDelete( NULL );
}
```


Creating a Task in FreeRTOS

```
portBASE_TYPE xTaskCreate( pdTASK_CODE pvTaskCode,
                           const signed char * const pcName,
                           unsigned short usStackDepth,
                           void *pvParameters,
                           unsigned portBASE_TYPE uxPriority,
                           xTaskHandle *pxCreatedTask
                           );
```

- pvTaskCode – A pointer to the task function
- pcName – Descriptive task name used during profiling
- usStackDepth – Stack depth/size
- pvParameters – Pointer to task parameters
- uxPriority – Task priority
- pxCreatedTask – Returns a task “handle”
- Returns success or failure

Embedded Real-Time Systems

17

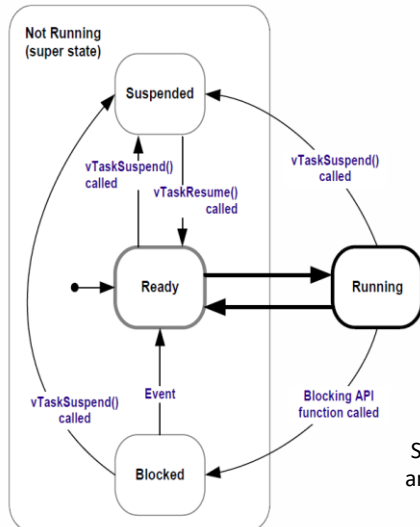
Major States of A Thread

- **Sleeping:** A thread is in sleeping state after creation or completion. It will stay in this state until it is released. It is not eligible for execution.
- **Ready:** A thread is in this state, when it released or preempted and is not blocked. It is eligible for execution.
- **Running:** A thread is in this state, if it is executing.
- **Suspended (Blocked):** A thread is in this state if it is released but has not yet completed, since it cannot proceed to a lack of a resource or budget.
- **Terminated:** A thread enters this state, if it will not execute again. Threads in this state may be destroyed.

Embedded Real-Time Systems

18

Full Task State Machine of FreeRTOS



Tasks can enter the “Blocked” state to wait for two types of events:

- Temporal Events, e.g. delaying for a fixed amount or until an absolute time
- Synchronization Events, e.g., waiting for data

Synchronization events can come in many forms and can be combined with delay events, e.g. wait for data until a time period has elapsed.

Embedded Real-Time Systems

19

System Calls

- A *system call* is a call to one of the RTOS API functions, like `xSemaphoreTake()` in the case of FreeRTOS.
- When the system call is executed, the **kernel saves the context** of the calling thread (and switches from user mode to **kernel mode**).
- The kernel executes then the system call and executes a **return from exception**(, which means that the system returns to user mode).
- Then the **ready thread with the highest priority** will be scheduled and starts to execute.

Embedded Real-Time Systems

20

Fixed-Priority Scheduling

- All modern operating systems support fixed-priority scheduling.
- The assigned priority is kept in the TCB.
- If protocols like priority inheritance protocol are used, even current priority has to be part of the TCB.
- Kernel maintains a ready queue for each priority level. When a thread is ready to execute, it is put in the queue that corresponds to its current priority.

Embedded Real-Time Systems

21

Dynamic Priorities

- Most operating systems support even dynamic priorities.
- This means that there is a system call that a thread can use in order to change its own or another threads priority.

Embedded Real-Time Systems

22

Example: Defining a Printer Task

```
void vTaskFunction( void *pvParameters )
{
    char *pcTaskName;

    /* The string to print out is passed in via the parameter. Cast this to a
    character pointer. */
    pcTaskName = ( char * ) pvParameters;

    /* As per most tasks, this task is implemented in an infinite loop. */
    for( ;; )
    {
        /* Print out the name of this task. */
        vPrintString( pcTaskName );

        /* Delay for a period. This time a call to vTaskDelay() is used which
        places the task into the Blocked state until the delay period has expired.
        The delay period is specified in 'ticks', but the constant
        portTICK_RATE_MS can be used to convert this to a more user friendly value
        in milliseconds. In this case a period of 250 milliseconds is being
        specified. */
        vTaskDelay( 250 / portTICK_RATE_MS );
    }
}
```

Creating the Tasks

```
/* Define the strings that will be passed in as the task parameters. These are
defined const and not on the stack to ensure they remain valid when the tasks are
executing. */
static const char *pcTextForTask1 = "Task 1 is running\n";
static const char *pcTextForTask2 = "Task 2 is running\n";

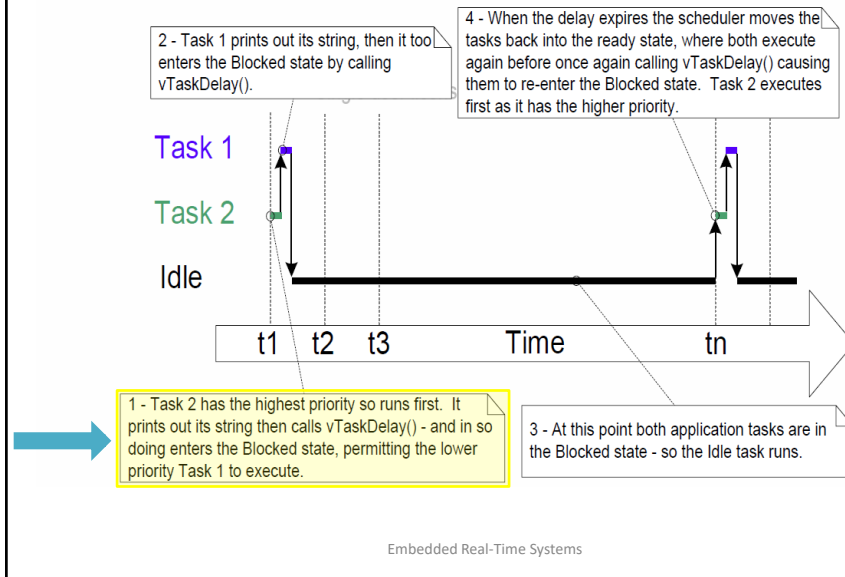
int main( void )
{
    /* Create the first task at priority 1. The priority is the second to last
    parameter. */
    xTaskCreate( vTaskFunction, "Task 1", 240, (void*)pcTextForTask1, 1, NULL );

    /* Create the second task at priority 2. */
    xTaskCreate( vTaskFunction, "Task 2", 240, (void*)pcTextForTask2, 2, NULL );

    /* Start the scheduler so the tasks start executing. */
    vTaskStartScheduler();

    /* If all is well we will never reach here as the scheduler will now be
    running. If we do reach here then it is likely that there was insufficient
    heap available for the Idle task to be created. */
    for( ;; );
}
```

Task Execution



Queue Access

- Accessed through a *handle* that is returned when the queue is created
- Operates as a **FIFO**, but can write to front
- Can also “peek”, i.e., receive a copy of data, or “query” the number of items in the queue
- Access can include a *maximum blocking time* to wait if the queue is full or empty (in ticks)
- Special API functions must be used to access from an Interrupt Service Routine (ISR)

Example using Queues

- Single queue – multiple send, single read task
- Send tasks are low priority and don't block
- Receive task is higher priority and does have a specified blocking time
- Thus, as soon as an item is written the read unblocks and empties queue
- As a result, queue should never have more than one item

Embedded Real-Time Systems

27

main()

```

xQueueHandle xQueue; // Global queue handle used by all tasks

int main( void )
{
    vSetupEnvironment(); // Configure HW

    /* The queue is created to hold a maximum of 5 long values. */
    xQueue = xQueueCreate( 5, sizeof( long ) );

    if( xQueue != NULL ) // Successfully created

        /* Two send tasks, priority 1. Each write 100 or 200 to queue */
        xTaskCreate( vSenderTask, ( signed char * ) "Sender1", 240, ( void * ) 100, 1, NULL );
        xTaskCreate( vSenderTask, ( signed char * ) "Sender2", 240, ( void * ) 200, 1, NULL );

        /* Read task, priority 2 */
        xTaskCreate( vReceiverTask, ( signed char * ) "Receiver", 240, NULL, 2, NULL );

        /* Start the scheduler so the created tasks start executing. */
        vTaskStartScheduler();
    }
    else
    {
        /* The queue could not be created. */
    }

    /* If all is well we will never reach here */
    for( ;; );
    return 0;
}

```

Embedded Real-Time Systems

28

Send Task

```
static void vSenderTask( void *pvParameters )
{
    long lValueToSend;
    portBASE_TYPE xStatus;

    /* Two instances will be created of this task */
    lValueToSend = ( long ) pvParameters;

    for( ;; )
    {
        xStatus = xQueueSendToBack( xQueue, &lValueToSend, 0 );
        /* The first parameter is the queue handle
        The second parameter is the address of the data to be sent.
        The third parameter is the Block time */

        if( xStatus != pdPASS )
        {
            /* We could not write to the queue it must be full */
            vPrintString( "Could not send to the queue.\r\n" );
        }

        /* Allow the other sender task to execute. */
        taskYIELD();
    }
}
```

Embedded Real-Time Systems

29

Receive Task

```
static void vReceiverTask( void *pvParameters )
{
    long lReceivedValue; // variable that holds received data
    portBASE_TYPE xStatus;
    const portTickType xTicksToWait = 100 / portTICK_RATE_MS;

    for( ;; )
    {
        if( uxQueueMessagesWaiting( xQueue ) != 0 )
            vPrintString( "Queue should have been empty!\r\n" );

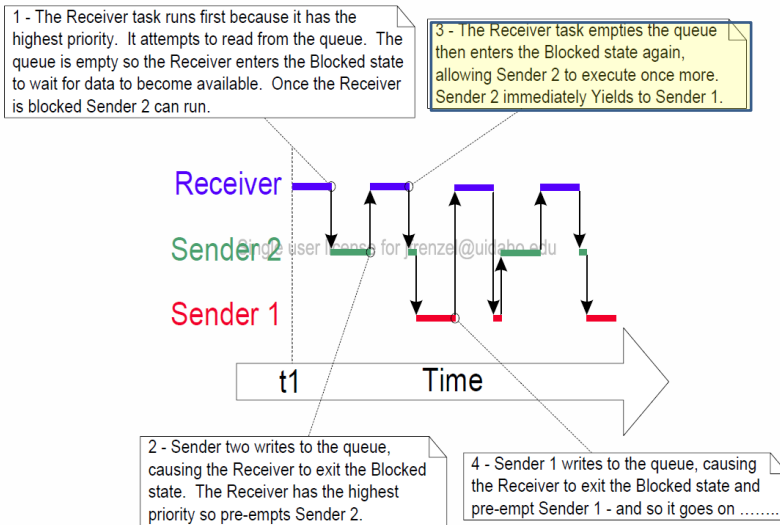
        xStatus = xQueueReceive( xQueue, &lReceivedValue, xTicksToWait );
        /* Queue handle, pointer to buffer, block time */

        if( xStatus == pdPASS )
        {
            vPrintStringAndNumber( "Received = ", lReceivedValue );
        }
        else
        {
            /* Did not receive anything, even after waiting for 100ms. */
            vPrintString( "Could not receive from the queue.\r\n" );
        }
    }
}
```

Embedded Real-Time Systems

30

Task Execution



Embedded Real-Time Systems

31

Tips

- For large amounts of data, send a pointer to the data (i.e., shared memory buffer)
- Care must be taken to ensure that sender doesn't modify RAM until it has been read
- If memory allocated dynamically, only one task should free the memory
- Beware sharing data on a task stack via a pointer, as the stack frame might change (Use global?)

Embedded Real-Time Systems

32

Binary Semaphores

- Can be used for synchronization between tasks
- Can also be used for connecting an ISR to a task
 - The “handler” task is written to block, waiting for a “binary semaphore” (i.e., signal)
 - The handler is trying to “take” a semaphore, and the ISR is written to “give” the semaphore
- A binary semaphore is analogous to a one-element queue, i.e., full or empty
- But rather than return to the preempted task, control is transferred to the handler ...

Example: Periodic Task

```
static void vPeriodicTask( void *pvParameters )
{
    /* As per most tasks, this task is implemented within an infinite loop. */
    for( ;; )
    {
        /* This task is just used to 'simulate' an interrupt. This is done by
        periodically generating a software interrupt. */
        vTaskDelay( 500 / portTICK_RATE_MS );

        /* Generate the interrupt, printing a message both beforehand and
        afterwards so the sequence of execution is evident from the output. */
        vPrintString( "Periodic task - About to generate an interrupt.\n" );
        mainTRIGGER_INTERRUPT();
        vPrintString( "Periodic task - Interrupt generated.\n\n" );
    }
}
```

Single user license for jfrenzel@uidaho.edu

Example: Periodic Task ISR

```
void vSW1_ISR_Handler( void )
{
    portBASE_TYPE xHigherPriorityTaskWoken = pdFALSE;

    /* 'Give' the semaphore to unblock the task. */
    xSemaphoreGiveFromISR( xBinarySemaphore, &xHigherPriorityTaskWoken );

    /* Clear the software interrupt flag. */
    mainCLEAR_INTERRUPT();

    /* Then clear the interrupt in the interrupt controller. */
    IFS0CLR = mainSW1_INT_BIT;

    /* Giving the semaphore may have unblocked a task - if it did and the
    unblocked task has a priority equal to or above the currently executing
    task then xHigherPriorityTaskWoken will have been set to pdTRUE and
    portEND_SWITCHING_ISR() will force a context switch to the newly unblocked
    higher priority task.

    NOTE: The syntax for forcing a context switch within an ISR varies between
    FreeRTOS ports. The portEND_SWITCHING_ISR() macro is provided as part of
    the PIC32 port layer for this purpose. taskYIELD() must never be called
    from an ISR! */
    portEND_SWITCHING_ISR( xHigherPriorityTaskWoken );
}
```

Embedded Real-Time Systems

35

Example: Periodic Task Task Handler

```
static void vHandlerTask( void *pvParameters )
{
    /* As per most tasks, this task is implemented within an infinite loop.

    Take the semaphore once to start with so the semaphore is empty before the
    infinite loop is entered. The semaphore was created before the scheduler
    was started so before this task ran for the first time.*/
    xSemaphoreTake( xBinarySemaphore, 0 );

    for( ;; )
    {
        /* Use the semaphore to wait for the event. The task blocks
        indefinitely meaning this function call will only return once the
        semaphore has been successfully obtained - so there is no need to check
        the returned value. */
        xSemaphoreTake( xBinarySemaphore, portMAX_DELAY );

        /* To get here the event must have occurred. Process the event (in this
        case we just print out a message). */
        vPrintString( "Handler task - Processing event.\n" );
    }
}
```

Embedded Real-Time Systems

36

Example: Periodic Task main ()

```
int main( void )
{
    /* Configure both the hardware and the debug interface. */
    vSetupEnvironment();

    /* Before a semaphore is used it must be explicitly created. In this example
    a binary semaphore is created. */
    vSemaphoreCreateBinary( xBinarySemaphore );

    /* Check the semaphore was created successfully. */
    if( xBinarySemaphore != NULL )
    {
        /* Create the 'handler' task. This is the task that will be synchronized
        with the interrupt. The handler task is created with a high priority to
        ensure it runs immediately after the interrupt exits. In this case a
        priority of 3 is chosen. */
        xTaskCreate( vHandlerTask, "Handler", 240, NULL, 3, NULL );

        /* Create the task that will periodically generate a software interrupt.
        This is created with a priority below the handler task to ensure it will
        get preempted each time the handler task exits the Blocked state. */
        xTaskCreate( vPeriodicTask, "Periodic", 240, NULL, 1, NULL );

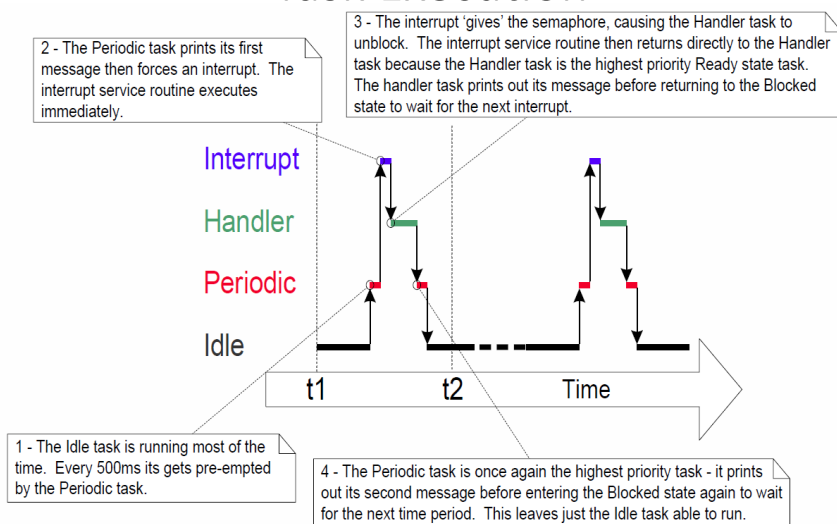
        /* Start the scheduler so the created tasks start executing. */
        vTaskStartScheduler();
    }

    /* If all is well we will never reach here as the scheduler will now be
    running the tasks. If we do reach here then it is likely that there was
    insufficient heap memory available for a resource to be created. */
    for( ;; );
}
```

Embedded Real-Time Systems

37

Example: Periodic Task Task Execution



Embedded Real-Time Systems

38

Some other FreeRTOS Facilities

- Counting semaphores
- Mutexes
- Software timers
- Event groups
- (Heap) memory management
- Dynamic priority management
- Kernel configuration

Embedded Real-Time Systems

39

RTOS Standards

- There exist a number of **standards** for real-time operating systems.
- At present the standards specify **portability at the source code level**, but the code needs to be recompiled for each operating system.
- At present there are the following main operating system standards
 - RT-POSIX - general-purpose operating system with real-time extensions
 - OSEK - automotive industry
 - APEX - avionic systems
 - μ ITRON - embedded systems

Embedded Real-Time Systems

40

RT-POSIX

- The real-time extension of POSIX, which is adopted by many operating systems.
- Specifies a large set of **system calls**, for **advanced RTOS functions**
 - priority inheritance
 - sporadic server scheduling
 - execution time budgeting
 - virtual memory management
- Four **real-time profiles** are defined by POSIX.13:
 - Minimal real-time system profile (PSE51) - small embedded systems with a memory footprint with tens of kilobyte.
 - Real-time controller profile (PSE52) - complete file system added
 - Dedicated real-time system profile (PSE53) - large embedded systems, includes memory protection
 - Multi-purpose real-time system profile (PSE54) - general purpose OSs

Examples:

- QNX
- RTEMS
- VxWorks