

Lecture 22: Quantitative Analysis

Seyed-Hosein Attarzadeh-Niaki

Based on slides by Edward Lee

Embedded Real-Time Systems

1

Review

- (Embedded) Real-time OSs
 - Characteristics
 - Microkernels
 - Tasks and scheduling
 - Queues and intercommunication
 - Semaphores and synchronization
 - Other facilities
- RTOS standards

Embedded Real-Time Systems

2


Outline

- Worst-case execution time problem
- Programs as Graphs
- Challenges of Execution Time Analysis
- Current Approaches
- Limitations and Future Directions

Embedded Real-Time Systems

3

Quantitative Analysis / Verification




Does the brake-by-wire software always actuate the brakes within 1ms?

Safety-critical embedded systems

Can this new app drain my iPhone battery in an hour?

Consumer devices



How much energy must the sensor node harvest for RSA encryption?

Energy-limited sensor nets, bio-medical apps, etc.

Embedded Real-Time Systems

4

Time is Central to Cyber-Physical Systems

Several timing analysis problems:

- Worst-case execution time (WCET) estimation
- Estimating the distribution of execution times
- Threshold property: can you produce a test case that causes a program to violate its deadline?
- Software-in-the-loop simulation: predict execution time of particular program path

➤ Various forms of the same basic problem.

Worst-Case Execution Time (WCET)

- The longest time taken by a software task to execute
 - Function of input data and environment conditions
- BCET = Best-Case Execution Time
 - shortest time taken by the task to execute

Consider this code: `*x = 10;`
on ARM Cortex-A9 MPCore dual core processor.

What's the WCET? BCET?

Worst-Case Execution Time (WCET) & BCET





Figure from R.Wilhelm et al., ACM Trans. Embed. Comput. Sys, 2007.

Embedded Real-Time Systems

7

The WCET Problem

- Given
 - the **code** for a software task
 - the **platform** (OS + hardware) that it will run on
- **Determine the WCET of the task.**



Embedded Real-Time Systems

8

The WCET Problem

- Why is this important? Where is the WCET used?

The WCET is central to the design of Embedded Systems:
 Needed for Correctness (does the task finish in time?) and
Performance (find optimal schedule for tasks)

- Can the WCET always be found?

In general, no, because the problem is *undecidable*.

- Is knowing WCET enough?

In general, no, because anomalies can occur when tasks end earlier than expected.

Typical WCET Problem

- Task executes within an infinite loop

```
while(1) {
    read_sensors();
    compute();
    write_to_actuators();
}
```

This code typically has:

- loops with finite bounds
- no recursion

Additional assumptions:

- runs uninterrupted
- single-threaded

Example Program: Modular Exponentiation

```

1 #define EXP_BITS 32
2
3 typedef unsigned int UI;
4
5 UI modexp(UI base, UI exponent, UI mod) {
6     int i;
7     UI result = 1;
8
9     i = EXP_BITS;
10    while(i > 0) {
11        if ((exponent & 1) == 1) {
12            result = (result * base) % mod;
13        }
14        exponent >>= 1;
15        base = (base * base) % mod;
16        i--;
17    }
18    return result;
19 }

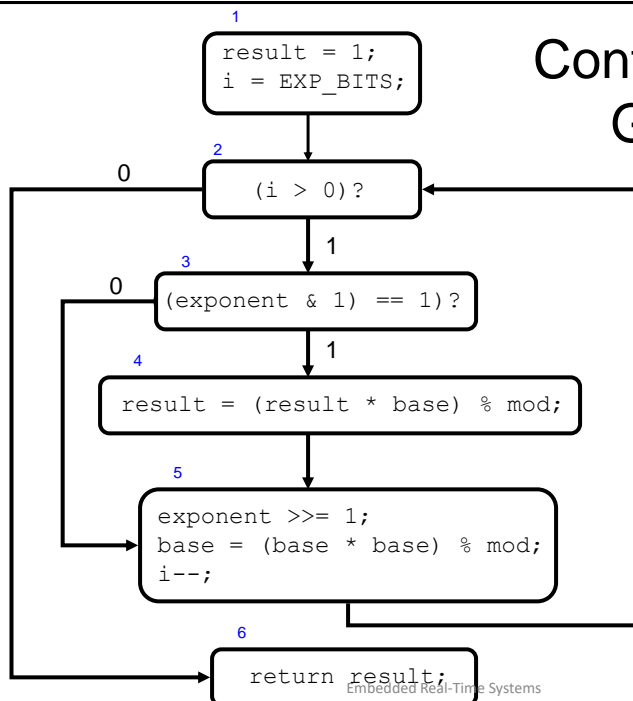
```

$$b^e = \begin{cases} (b^2)^{e/2} = (b^{e/2})^2, & e \text{ is even,} \\ (b^2)^{(e-1)/2} \cdot b = (b^{(e-1)/2})^2 \cdot b, & e \text{ is odd.} \end{cases}$$

Embedded Real-Time Systems

11

Control-Flow Graph



Each node is a
basic block

Embedded Real-Time Systems

12

Components of Execution Time Analysis (traditional approaches)

- **Program path (Control flow) analysis**
 - Want to find longest path through the program
 - Find loop bounds
 - Identify feasible paths through the program
 - Identify dependencies amongst different code fragments
- **Processor behavior analysis**
 - For small code fragments (basic blocks), generate bounds on run-times on the platform
 - Model details of architecture, including cache behavior, pipeline stalls, branch prediction, etc.
- **Outputs of both analyses feed into each other**

Embedded Real-Time Systems

13

Program Path Analysis: Path Explosion

```

/* MAXSIZE = 100 */
for (Outer = 0; Outer < MAXSIZE; Outer++) {
    for (Inner = 0; Inner < MAXSIZE; Inner++) {
        if (Array[Outer][Inner] >= 0) {
            Ptotal += Array[Outer][Inner];
            Pcnt++;
        } else {
            Ntotal += Array[Outer][Inner];
            Ncnt++;
        }
    }
    Postotal = Ptotal;
    Poscnt = Pcnt;
    Negtotal = Ntotal;
    Negcnt = Ncnt;
}

```

How many paths?

 2^{10000}

Example cnt.c from WCET benchmarks, Mälardalen Univ.

Embedded Real-Time Systems

14

Program Path Analysis: Determining Loop Bounds

```

1  #define EXP_BITS 32
2
3  typedef unsigned int UI;
4
5  UI modexp(UI base, UI exponent, UI mod) {
6      int i;
7      UI result = 1;
8
9      i = EXP_BITS;
10     while(i > 0) {
11         if ((exponent & 1) == 1) {
12             result = (result * base) % mod;
13         }
14         exponent >>= 1;
15         base = (base * base) % mod;
16         i--;
17     }
18     return result;
19 }

```

How many times around the while loop?

32

Embedded Real-Time Systems

15

Common Current Approach (high-level)

1. Manually construct processor behavior model
2. Use model to find “worst-case” starting processor states for each basic block → measure execution times of the blocks from these states
3. Use these times as upper bounds on the time of each basic block
4. Formulate an integer linear program to find the maximum sum of these bounds along any program path

Embedded Real-Time Systems

16

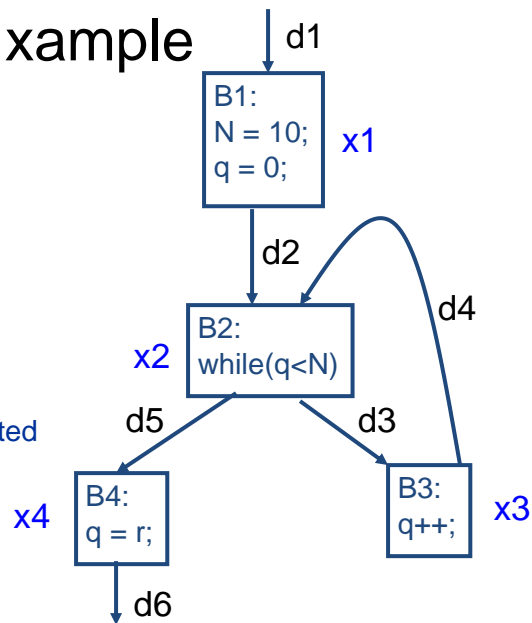
Example

```

N = 10;
q = 0;
while(q < N)
    q++;
q = r;

```

$x_i \rightarrow$ # times B_i is executed
 $d_j \rightarrow$ # times edge is executed



Example due to Y.T. Li and S. Malik

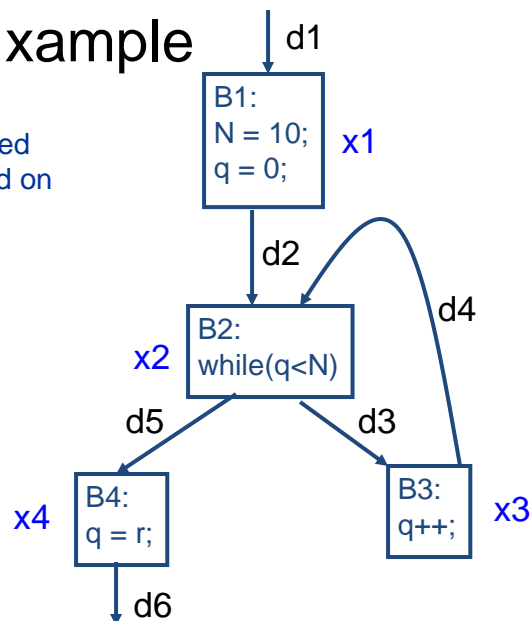
Embedded Real-Time Systems

17

Example

$x_i \rightarrow$ # times B_i is executed
 $d_j \rightarrow$ # times edge is executed
 $C_i \rightarrow$ measured upper bound on time taken by B_i

Want to
 maximize $\sum_i C_i x_i$
 subject to constraints
 $x_1 = d_1 = d_2$
 $d_1 = 1$
 $x_2 = d_2 + d_4 = d_3 + d_5$
 $x_3 = d_3 = d_4 = 10$
 $x_4 = d_5 = d_6$



Example due to Y.T. Li and S. Malik

Embedded Real-Time Systems

18

Integer Linear Programming

$x_i \rightarrow$ # times B_i is executed
 $d_j \rightarrow$ # times edge is executed
 $C_i \rightarrow$ measured upper bound on time taken by B_i

Want to
 maximize $\sum_i C_i x_i$
 subject to constraints
 $x_1 = d_1 = d_2$
 $d_1 = 1$
 $x_2 = d_2 + d_4 = d_3 + d_5$
 $x_3 = d_3 = d_4 = 10$
 $x_4 = d_5 = d_6$

The problem to solve is:

Find integer values x_i that maximize the sum (total execution time) subject to the (linear) constraint equations.

In general, this is an **Integer Linear Programming** (ILP) problem, which is known to be NP-hard, but for which there is good software.

Embedded Real-Time Systems

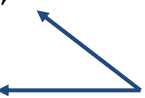
19

Program Path Analysis: Dependencies

```

void altitude_pid_run(void) {
    float err = estimator_z - desired_altitude;
    desired_climb = pre_climb + altitude_pgain * err;
    if (desired_climb < -CLIMB_MAX)
        desired_climb = -CLIMB_MAX;
    if (desired_climb > CLIMB_MAX)
        desired_climb = CLIMB_MAX;
}

```



Only one of these statements is executed
(CLIMB_MAX = 1.0)

How many *feasible* paths?

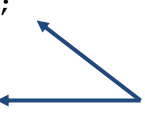
Example from "PapaBench" UAV autopilot code, IRIT, France

Embedded Real-Time Systems

20

Program Path Analysis: Dependencies

```
void altitude_pid_run(void) {
    float err = estimator_z - desired_altitude;
    desired_climb = pre_climb + altitude_pgain * err;
    if (desired_climb < -CLIMB_MAX)
        desired_climb = -CLIMB_MAX;
    if (desired_climb > CLIMB_MAX)
        desired_climb = CLIMB_MAX;
}
```



Only one of these statements is executed
(CLIMB_MAX = 1.0)

This adds an additional constraint of
the form:
 $x_i + x_j \leq 1$

Example from "PapaBench" UAV autopilot code, IRIT, France

Embedded Real-Time Systems

21

But the problem gets even harder...

Timing Analysis and Compositionality

Consider a task T with two parts A and B
composed in sequence: $T = A; B$

Is $WCET(T) = WCET(A) + WCET(B)$?

NOT ALWAYS!

WCETs cannot simply be composed ☹️

→ Due to dependencies "through environment"

Embedded Real-Time Systems

22

Processor Behavior Analysis: Cache Effects

```

1 float dot_product(float *x, float *y, int n) {
2     float result = 0.0;
3     int i;
4     for(i=0; i < n; i++) {
5         result += x[i] * y[i];
6     }
7     return result;
8 }

```

Suppose:

1. 32-bit processor
2. Direct-mapped cache holds two sets
 - 4 floats per set
 - x and y stored contiguously starting at address 0x0

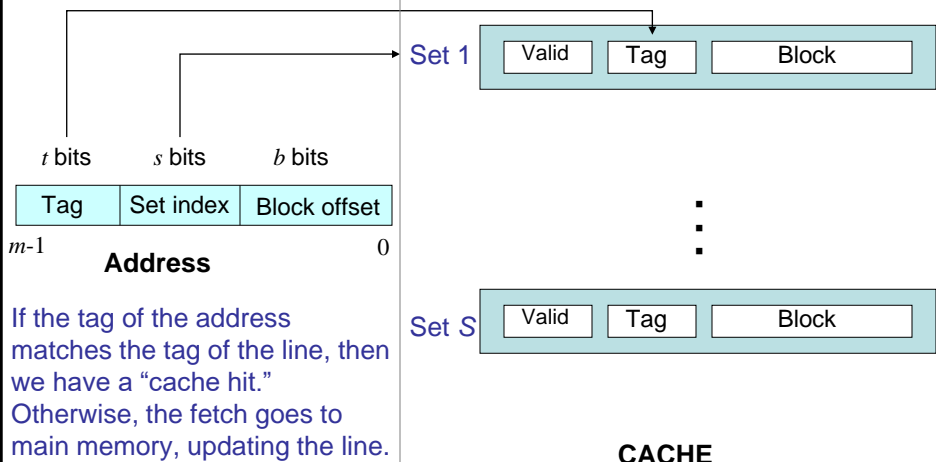
What happens when $n=2$?

Embedded Real-Time Systems

23

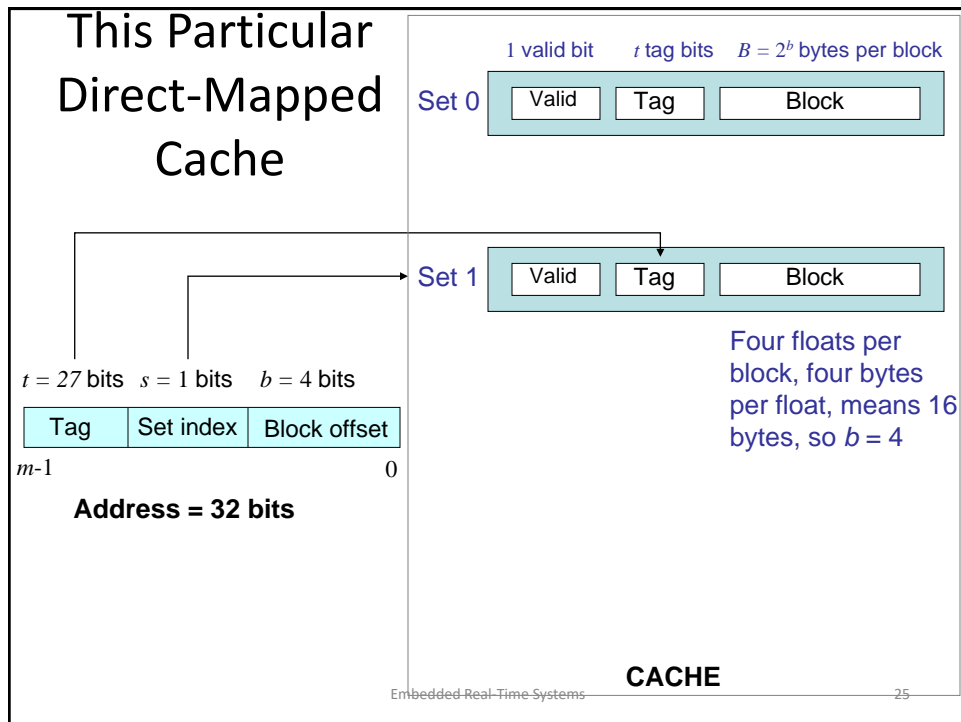
Recall Direct-Mapped Cache

A "set" consists of one "line"



Embedded Real-Time Systems

24



Processor Behavior Analysis: Cache Effects

```

1 float dot_product(float *x, float *y, int n) {
2     float result = 0.0;
3     int i;
4     for(i=0; i < n; i++) {
5         result += x[i] * y[i];
6     }
7     return result;
8 }

```

What happens when $n=2$?

Suppose:

1. 32-bit processor
2. Direct-mapped cache holds two sets
 - 4 floats per set
 - x and y stored contiguously starting at address $0x0$

$x[0]$ will miss, pulling $x[0]$, $x[1]$, $y[0]$ and $y[0]$ into the set 0. All but one access will be a cache hit.

Processor Behavior Analysis: Cache Effects

```

1 float dot_product(float *x, float *y, int n) {
2     float result = 0.0;
3     int i;
4     for(i=0; i < n; i++) {
5         result += x[i] * y[i];
6     }
7     return result;
8 }

```

What happens
when **n=8**?

*x[0] will miss,
pulling x[0-3] into
the set 0. Then
y[0] will miss,
pulling y[0-3] into
the same set,
evicting x[0-3].
Every access will
be a miss!*


Suppose:

- 32-bit processor
- Direct-mapped cache holds two sets
 - 4 floats per set
 - x and y stored contiguously starting at address 0x0

Embedded Real-Time Systems

27

Timing Anomalies



Scenario 1: Block A hits in I-cache, triggers branch speculation, and prefetch of instructions, then predicted branch is wrong, so Block B must execute, but it's been evicted from I-cache, execution of B delayed.

Scenario 2: Block A misses in I-cache, no branch prediction, then B hits in I-cache, B completes.

[from R.Wilhelm et al., ACM Trans. Embed. Comput. Sys, 2007.]

Embedded Real-Time Systems

28

How to Measure Run-Time

Several techniques, with varying accuracy

- Instrument code to sample CPU **cycle counter**
 - relatively easy to do, read processor documentation for assembly instruction
- Use **cycle-accurate simulator** for processor
 - useful when hardware is not available/ready
- Use **Logic Analyzer**
 - non-intrusive measurement, more accurate
- ...

Embedded Real-Time Systems

29

Measurement Pitfalls

- Instrumentation incurs small **overhead**
 - measure long enough code sequence to compensate
- **Multi-tasking effects**: counter keeps going even when the task of interest is inactive
 - take multiple measurements and pick “k best” (cluster)
- **Multicores/hyperthreading**
 - Need to ensure that task is ‘locked’ to a single core
- **Power management effects**
 - CPU speed might change, timer could get reset during hibernation

Embedded Real-Time Systems

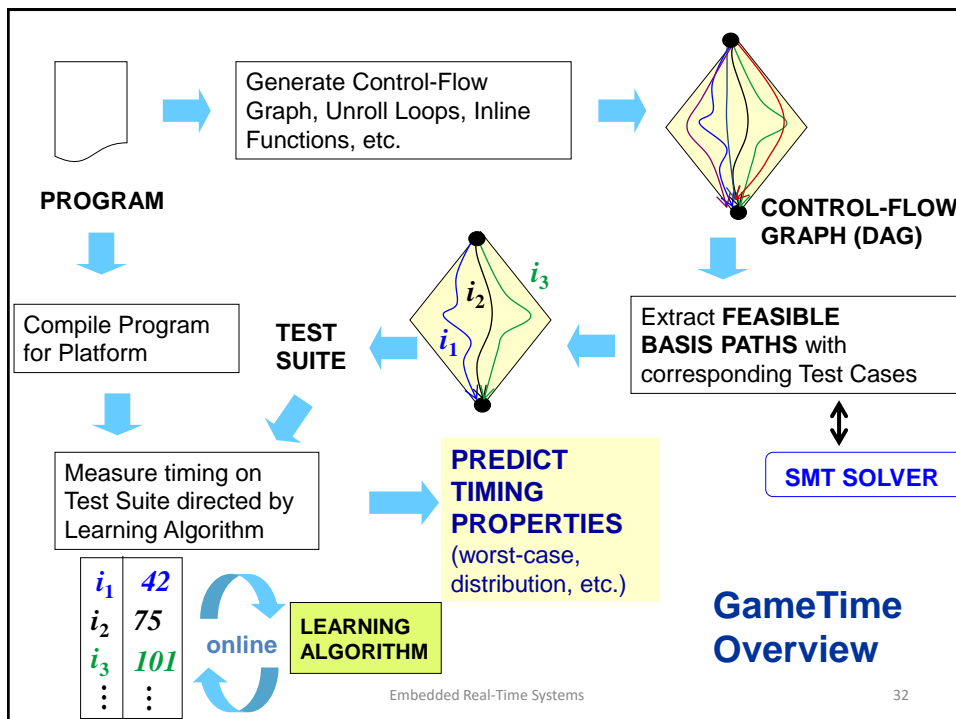
30

Some WCET Estimation Tools

- Commercial Tools: aiT, RapiTime, ...
- University/Research Tools: GameTime, Chronos, ...
 - GameTime: combines machine learning with automated theorem proving technology
 - Uses Programs as Graphs and integer linear programming ideas, but almost everything else different from other tools!

Embedded Real-Time Systems

31



Embedded Real-Time Systems

32

Open Problems

- **Architectures** are getting much **more complex**.
 - Can we create processor models without the agonizing pain?
[Yes, employ machine learning from systematic measurements
→ GameTime project]
 - Can we change the architecture to make timing analysis easier?
[Yes, see PRET machine project]
- **Analysis methods are “Brittle”** – small changes to code and/or architecture can require completely re-doing the WCET computation
 - GameTime project addresses this:
 - Use robust techniques that learn about processor/platform behavior
 - Need to deal with concurrency, e.g., interrupts
- Need more **reliable ways to measure execution time**

Embedded Real-Time Systems

33

PRET Machines – Giving Software the Capabilities their Hardware Already Has.

- PREcision-Timed processors = PRET
 - Predictable, REpeatable Timing = PRET
 - Performance with REpeatable Timing = PRET
- <http://chess.eecs.berkeley.edu/pret>

```
// Perform the convolution.
for (int i=0; i<10; i++) {
    x[i] = a[i]*b[j-i];
    // Notify listeners.
    notify(x[i]);
}
```

Computing

+



= PRET

With time

Embedded Real-Time Systems

34

Major Challenges and existence proofs that they can be met


- Pipelines
 - fine-grain multithreading
- Memory hierarchy
 - memory controllers with controllable latency
- I/O
 - threaded interrupts, with bounded effects on timing

Embedded Real-Time Systems

35

Second Generation PRET

PTArm, a soft core on a
Xilinx Virtex 5 FPGA (2012)



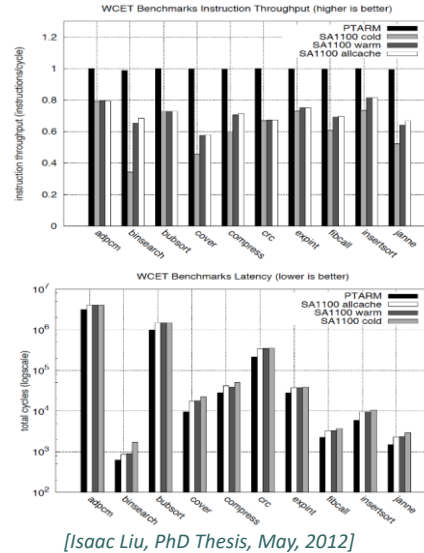
Embedded Real-Time Systems

36

Performance Cost?

No!

- The PRET project has shown that you do not need to sacrifice performance to get control over timing.
- But in PtARM, you need enough concurrency to not lose performance.

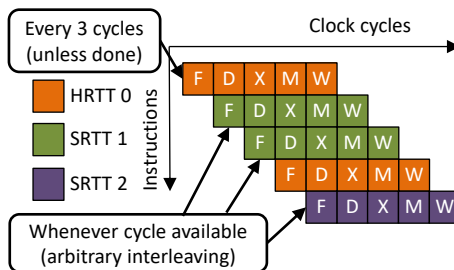


Embedded Real-Time Systems

37

Third-Generation PRET: Open-Source FlexPRET (Zimmer 2014/15)

- 32-bit, 5-stage thread interleaved pipeline, RISC-V ISA
 - Hard real-time HW threads:** scheduled at constant rate for isolation and repeatability.
 - Soft real-time HW threads:** share all available cycles for efficiency.
- Deployed on Xilinx FPGA (area comparable to Microblaze)



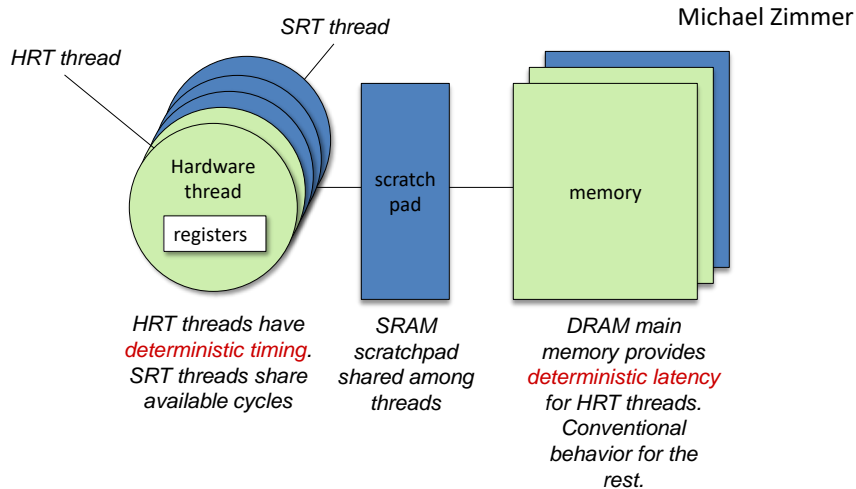
Digilent Atlys (Spartan 6) and NI myRIO (Zync)

Embedded Real-Time Systems

38


FlexPRET

*Hard-Real-Time (HRT) Threads
Interleaved with Soft-Real-Time (SRT) Threads*



FlexPRET I/O

Interrupt-Driven I/O is notorious for disrupting timing



FlexPRET Shows

- Not only is there no performance cost for appropriate workloads, but there is also no performance cost for inappropriate workloads!
- Pipelining, memory hierarchy, and interrupt-driven I/O can all be done without losing timing determinacy!

Industry has shown some interest

ARM The Architecture for the Digital World®

Products Support Community Markets About Careers

You are here: > Products > Processors > Cortex-R Series > Cortex-R52 Processor

Cortex-R Series

- > **Cortex-R52 Processor**
- > Cortex-R8 Processor
- > Cortex-R7 Processor
- > Cortex-R5 Processor
- > Cortex-R4 Processor

Cortex-R52 Processor

The ARM® Cortex®-R52 delivers the highest level of integrated capability for functional safety of any ARM processor. It builds on the capabilities of the Cortex-R5 processor to meet the rising performance needs of advanced real-time embedded systems.

Cortex-R52 has been designed specifically to address the needs of functional safety. It has been developed within a process to fulfil IEC 61508 SIL 3 and ISO 26262 ASIL D requirements enabling it to be readily adopted for highly dependable systems where safety is critical.

Cortex-R52 implements hardware to simplify the integration of increasingly complex real-time software environments while providing the robust separation of software necessary to protect safety-critical code. As the first ARMv8-R processor, Cortex-R52 introduces an extra privilege level which provides support for a hypervisor.

This is all achieved without impacting the determinism needed for real time systems and while providing higher levels of performance from single and multicore configurations.

Key Benefits

- Streamlined next generation system development with the highest level of integrated functional safety features
- Simplified integration of complex software through strong separation of mixed criticality code without impacting real-time performance
- Address higher workloads with increased performance compared to Cortex-R5

ARM Cortex-R52

ARMv8-A CPU, L2P, MPU, NEON, GIC, SCU, PPU, TCM, AXI, ETH, Flash, AOS

[Request More Information](#)

Optional Reading Material

Embedded Real-Time Systems

43

Cycle Counters

- Most modern systems have built in registers that are incremented every clock cycle
- Special assembly code instruction to access
- On Intel 32-bit x86 machines since Pentium:
 - 64 bit counter
 - RDTSC instruction (Read Time Stamp Counter) sets %edx register to high order 32-bits, %eax register to low order 32-bits
- Wrap-around time for 2 GHz machine
 - Low order 32-bits every 2.1 seconds
 - High order 64 bits every 293 years

[slide due to R. E. Bryant and D. R. O'Hallaron]

Embedded Real-Time Systems

44

Measuring with Cycle Counter

- Idea
 - Get current value of cycle counter
 - store as pair of unsigned's cyc_hi and cyc_lo
 - Compute something
 - Get new value of cycle counter
 - Perform double precision subtraction to get elapsed cycles

```
/* Keep track of most recent reading of cycle counter */
static unsigned cyc_hi = 0;
static unsigned cyc_lo = 0;

void start_counter()
{
    /* Get current value of cycle counter */
    access_counter(&cyc_hi, &cyc_lo);
}
```

[slide due to R. E. Bryant and D. R. O'Hallaron]

Embedded Real-Time Systems

45

Accessing the Cycle Counter

- GCC allows inline assembly code with mechanism for matching registers with program variables
- Code only works on x86 machine compiling with GCC

```
void access_counter(unsigned *hi, unsigned *lo)
{
    /* Get cycle counter */
    asm("rdtsc; movl %%edx,%0; movl %%eax,%1"
        : "=r" (*hi), "=r" (*lo)
        : /* No input */
        : "%edx", "%eax");
}
```

- Emit assembly with rdtsc and two movl instructions

[slide due to R. E. Bryant and D. R. O'Hallaron]

Embedded Real-Time Systems

46

Completing Measurement

- Get new value of cycle counter
- Perform double precision subtraction to get elapsed cycles
- Express as `double` to avoid overflow problems

```
double get_counter()
{
    unsigned ncyc_hi, ncyc_lo
    unsigned hi, lo, borrow;
    /* Get cycle counter */
    access_counter(&ncyc_hi, &ncyc_lo);
    /* Do double precision subtraction */
    lo = ncyc_lo - cyc_lo;
    borrow = lo > ncyc_lo;
    hi = ncyc_hi - cyc_hi - borrow;
    return (double) hi * (1 << 30) * 4 + lo;
}
```

[slide due to R. E. Bryant and D. R. O'Hallaron]

47

Timing With Cycle Counter

- Time Function P
 - First attempt: Simply count cycles for one execution of P

```
double tcycles;
start_counter();
P();
tcycles = get_counter();
```

- What can go wrong here?

[slide due to R. E. Bryant and D. R. O'Hallaron]

48

Dealing with Overhead & Cache Effects

- Always execute function once to “warm up” cache
- Keep doubling number of times execute P() until reach some threshold
 - Used CMIN = 50000

```
int cnt = 1;
double cmeas = 0;
double cycles;
do {
    int c = cnt;
    P(); /* Warm up cache */
    get_counter();
    while (c-- > 0)
        P();
    cmeas = get_counter();
    cycles = cmeas / cnt;
    cnt += cnt;
} while (cmeas < CMIN); /* Make sure have enough */
return cycles / (1e6 * MHZ);
```

Embedded Real-Time Systems

49

Timing With Cycle Counter

- Determine Clock Rate of Processor
 - Count number of cycles required for some fixed number of seconds

```
double MHZ;
int sleep_time = 10;
start_counter();
sleep(sleep_time);
MHZ = get_counter() / (sleep_time * 1e6);
```

- Time Function P
 - First attempt: Simply count cycles for one execution of P

```
double tsecs;
start_counter();
P();
tsecs = get_counter() / (MHZ * 1e6);
```

[slide due to R. E. Bryant and D. R. O'Hallaron]

Embedded Real-Time Systems

50

Measurement Pitfalls

- Instrumentation incurs small overhead
 - measure long enough code sequence to compensate
- Cache effects can skew measurements
 - “warm up” the cache before making measurement
- Multi-tasking effects: counter keeps going even when the task of interest is inactive
 - take multiple measurements and pick “k best” (cluster)
- Multicores/hyperthreading
 - Need to ensure that task is ‘locked’ to a single core
- Power management effects
 - CPU speed might change, timer could get reset during hibernation